

1. American Sign Language (ASL)

Import packages and set numpy random seed

```
import numpy as np
np.random.seed(5)
import tensorflow as tf
tf.set_random_seed(2)
from datasets import sign_language
import matplotlib.pyplot as plt
%matplotlib inline
```

Load pre-shuffled training and test datasets

```
(x_train, y_train), (x_test, y_test) = sign_language.load_data()
```

Using TensorFlow backend.


2. Visualize the training data

Store labels of dataset

```
labels = ['A', 'B', 'C']
```

Print the first several training images, along with the labels

```
fig = plt.figure(figsize=(20,5))
for i in range(36):
    ax = fig.add_subplot(3, 12, i + 1, xticks=[], yticks=[])
    ax.imshow(np.squeeze(x_train[i]))
    ax.set_title("{} {}".format(labels[y_train[i]]))
plt.show()
```



3. Examine the dataset

Number of A's in the training dataset

```
num_A_train = sum(y_train==0)
```

Number of B's in the training dataset

```
num_B_train = sum(y_train==1)
```

Number of C's in the training dataset

```
num_C_train = sum(y_train==2)
```

Number of A's in the test dataset

```
num_A_test = sum(y_test==0)
```

Number of B's in the test dataset

```
num_B_test = sum(y_test==1)
```

```

# Number of C's in the test dataset
num_C_test = sum(y_test==2)

# Print statistics about the dataset
print("Training set:")
print("\tA: {}, B: {}, C: {}".format(num_A_train, num_B_train,
num_C_train))
print("Test set:")
print("\tA: {}, B: {}, C: {}".format(num_A_test, num_B_test,
num_C_test))

```

```

Training set:
    A: 540, B: 528, C: 532
Test set:
    A: 118, B: 144, C: 138

```

4. One-hot encode the data

```

from keras.utils import np_utils

# One-hot encode the training labels
y_train_OH = np_utils.to_categorical(y_train)

# One-hot encode the test labels
y_test_OH = np_utils.to_categorical(y_test)

```

5. Define the model

```

from keras.layers import Conv2D, MaxPooling2D
from keras.layers import Flatten, Dense
from keras.models import Sequential

model = Sequential()
# First convolutional layer accepts image input
model.add(Conv2D(filters=5, kernel_size=5, padding='same',
activation='relu',
                    input_shape=(50, 50, 3)))
# Add a max pooling layer
model.add(MaxPooling2D((4, 4)))
# Add a convolutional layer
model.add(Conv2D(filters=15, kernel_size=5, padding='same',
activation='relu'))
# Add another max pooling layer
model.add(MaxPooling2D((4, 4)))
# Flatten and feed to output layer
model.add(Flatten())
model.add(Dense(3, activation='softmax'))

# Summarize the model
model.summary()

```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 50, 50, 5)	380
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 5)	0
conv2d_2 (Conv2D)	(None, 12, 12, 15)	1890
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 15)	0
flatten_1 (Flatten)	(None, 135)	0
dense_1 (Dense)	(None, 3)	408
Total params: 2,678		
Trainable params: 2,678		
Non-trainable params: 0		

6. Compile the model

Compile the model

```
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

7. Train the model

Train the model

```
hist = model.fit(x=x_train, y=y_train_OH, epochs=2,
                 validation_split=0.2, batch_size=32)
```

Train on 1280 samples, validate on 320 samples

Epoch 1/2

1280/1280 [=====] - 2s 2ms/step - loss: 0.9560 - acc: 0.6164 - val_loss: 0.7627 - val_acc: 0.8500

Epoch 2/2

1280/1280 [=====] - 2s 2ms/step - loss: 0.6086 - acc: 0.8875 - val_loss: 0.4636 - val_acc: 0.9187

8. Test the model

Obtain accuracy on test set

```
score = model.evaluate(x=x_test,
                       y=y_test_OH,
                       verbose=0)
print('Test accuracy:', score[1])
```

Test accuracy: 0.94

9. Visualize mistakes

```
# Get predicted probabilities for test dataset
```

```
y_probs = model.predict(x_test)
```

```
# Get predicted labels for test dataset
```

```
y_preds = np.argmax(y_probs, axis=1)
```

```
# Indices corresponding to test images which were mislabeled
```

```
bad_test_idx = np.arange(len(y_test))[y_test != y_preds]
```

```
# Print mislabeled examples
```

```
fig = plt.figure(figsize=(25,4))
```

```
for i, idx in enumerate(bad_test_idx):
```

```
    ax = fig.add_subplot(2, np.ceil(len(bad_test_idx)/2), i + 1,  
xticks=[], yticks=[])
```

```
    ax.imshow(np.squeeze(x_test[idx]))
```

```
    ax.set_title("{} (pred: {})".format(labels[y_test[idx]],  
labels[y_preds[idx]]))
```

